

COMP90038

Algorithms and Complexity

Lecture 15: Balanced Trees

(with thanks to Harald Søndergaard & Michael Kirley)

Casey Myers

Casey.Myers@unimelb.edu.au

David Caro Building (Physics) 274


Review from Lecture 14

- A **binary search tree**, or **BST**, is a binary tree that stores elements in all internal nodes, with each sub-tree satisfying the BST property:

Let the root be r ; then each element in the left subtree is smaller than r and each element in the right sub-tree is larger than r . (For simplicity, we assume that all keys are different.)


Review from Lecture 14

- BSTs are useful for search applications. To search for k in a BST, compare against its root r . If $r = k$, we are done; otherwise search in the left or right sub-tree, according to $k < r$ or $k > r$.
- If a BST with n elements is “reasonably” balanced, search involves in the worst case, $\Theta(\log n)$ comparisons.



Review from Lecture 14

- If the BST is not well balanced, search performance degrades, and may be as bad as linear search:



Approaches to Balanced Binary Search Trees

- To optimise the performance of BST search, it is important to keep trees (reasonably) balanced
- Instance simplification approaches: Self-balancing trees:
 - AVL trees
 - Red-black trees
 - Splay trees
- Representational changes:
 - 2-3 trees
 - 2-3-4 trees
 - B-trees

Approaches to Balanced Binary Search Trees





- To optimise the performance of BST search, it is important to keep trees (reasonably) balanced
- Instance simplification approaches: Self-balancing trees:
 - AVL trees
 - Red-black trees
 - Splay trees
- Representational changes:
 - 2-3 trees
 - 2-3-4 trees
 - B-trees

AVL Trees

- Named after Adelson-Velsky and Landis
- Recall that we defined the height of the empty tree as -1.
- For a binary (sub-) tree, let the **balance factor** be the difference between the height of its left sub-tree and that of its right sub-tree.
- An **AVL tree** is a BST in which the balance factor is -1, 0, or 1, for every sub-tree.





AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?





AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?





AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?







H: 3, 1; BF: $3 - 1 = 2$



AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?





H: 3, 3; BF: 3-3 = 0



AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?





H: 0, 1; BF: 0-1 = -1



AVL Trees: Examples and Counter-Examples

- Which of these are AVL trees?


H: 0, 2; BF: $0-2 = -2$




Building an AVL Tree

- As with standard BSTs, insertion of a new node always takes place at the fringe of the tree.
- If insertion of the new node makes the AVL tree unbalanced (some nodes get balance factors of 2 or -2), transform the tree to regain its balance.
- Regaining balance can be achieved with one or two simple, local transformations, so-called **rotations**.


AVL Trees: R-Rotation




AVL Trees: L-Rotation



AVL Trees: LR-Rotation




AVL Trees: RL-Rotation



AVL Trees: Where to Perform the Rotation


- Along an unbalanced path, we may have several nodes with balance factor 2 (or -2).



- It is always the **lowest** unbalanced subtree that is re-balanced.

AVL Trees: Where to Perform the Rotation


- Along an unbalanced path, we may have several nodes with balance factor 2 (or -2).



- It is always the **lowest** unbalanced subtree that is re-balanced.

AVL Trees: Where to Perform the Rotation

- Along an unbalanced path, we may have several nodes with balance factor 2 (or -2).




- It is always the **lowest** unbalanced subtree that is re-balanced.

AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12


AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 14, 17, 11, 7, 53, 4, 13, 12




AVL Trees: The Single Rotation, Generally



- This shows an **R-rotation**; an **L-rotation** is similar.

AVL Trees: The Single Rotation, Generally




- This shows an **LR-rotation**; an **RL-rotation** is similar.

AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7


AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7



AVL Trees: Example


- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7



AVL Trees: Example


- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7

H: 3, 1; BF: $3-1 = 2$






AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




\Rightarrow




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7



AVL Trees: Example

- Insert each of these number into the AVL tree in order: 5, 6, 8, 3, 2, 4, 7




Properties of AVL Trees

- The notion of “balance” that is implied by the AVL condition is sufficient to guarantee that the depth of an AVL tree with n nodes is $\Theta(\log n)$.
- For random data, the depth is very close to $\log_2 n$, the optimum.
- In the worst case, search will need at most 45% more comparisons than with a perfectly balanced BST.
- **Deletion** is harder to implement than insertion but also $\Theta(\log n)$.

Other Kinds of Balanced Trees

- A **red-black** tree is a BST with a slightly different concept of “balanced”. Its nodes are coloured red or black, so that:
 - No red node has a red child.
 - Every path from the root to the fringe has the same number of black nodes.



A worst-case red-black tree
(the longest path is twice as
long as the shortest path).


A **splay** tree is a BST which is not only self-adjusting, but also **adaptive**. Frequently accessed items are brought closer to the root, so their access becomes cheaper.

2-3 Trees


- 2-3 trees and 2-3-4 trees are search trees that allow more than one item to be stored in a tree node.
- As with BSTs, a node that holds a single item has (at most) two children.
- A node that holds two items (a so-called **3-node**) has (at most) three children.
- And for 2-3-4 trees, a node that holds three items (a **4-node**) has (at most) four children.
- This allows for a simple way of keeping search trees balanced.

2-Nodes and 3-Nodes

2-node



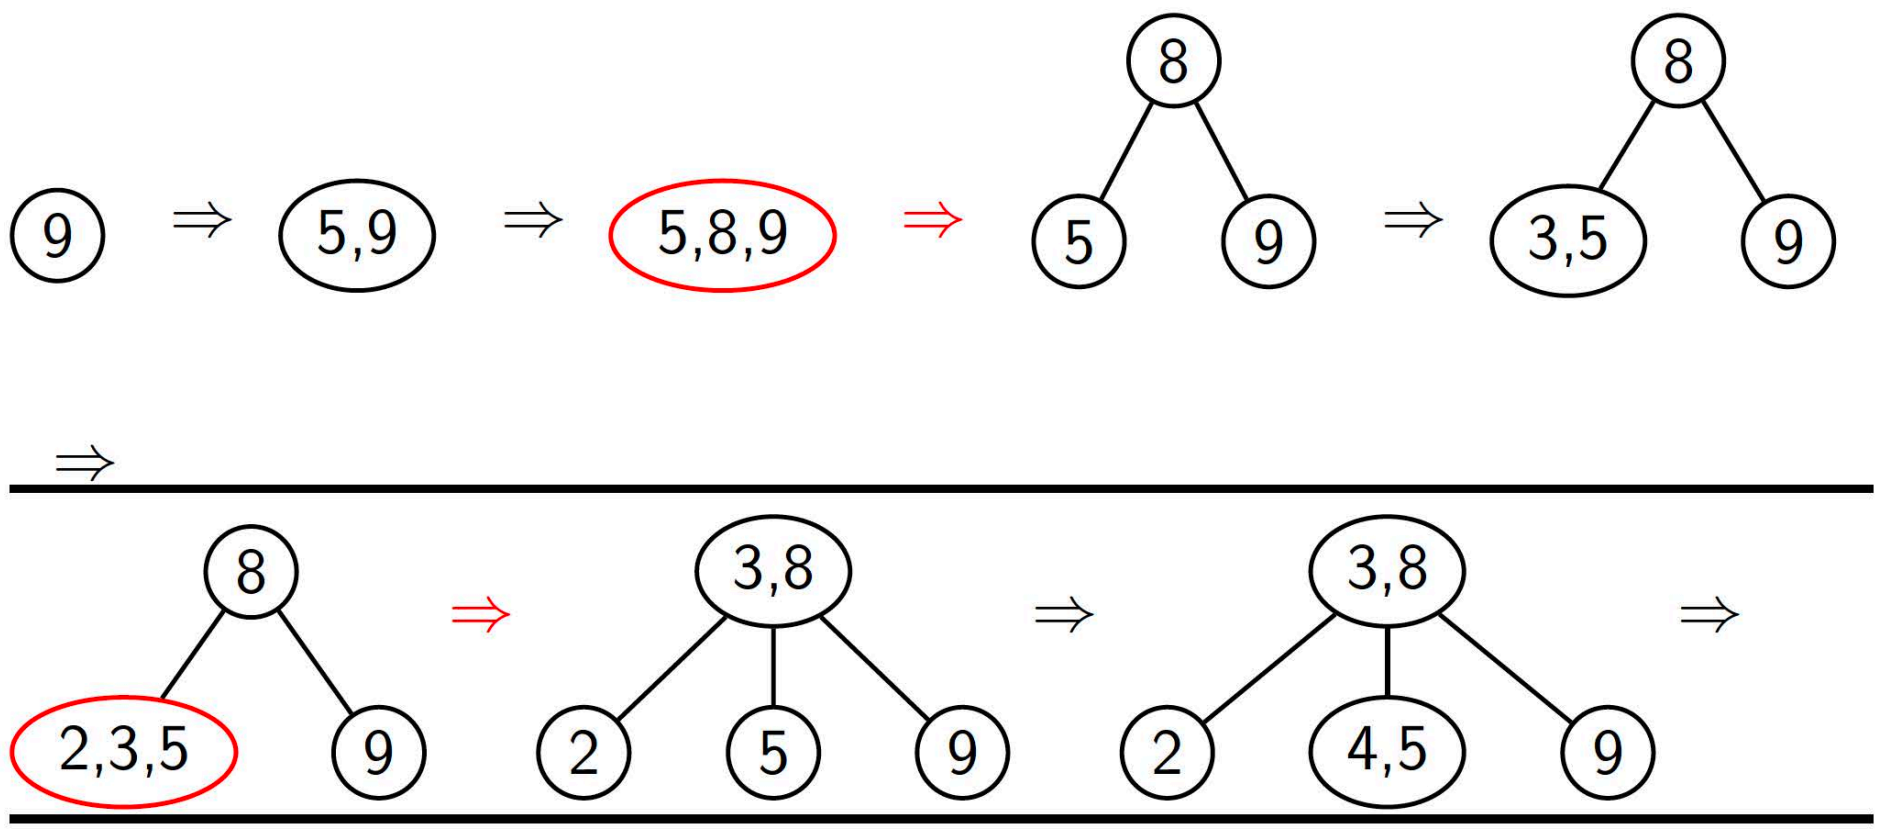
3-node




Insertion in a 2-3 Tree

- To insert a key k , pretend that we are searching for k .
- This will take us to a leaf node in the tree, where k should now be inserted; if the node we find there is a 2-node, k can be inserted without further ado.
- Otherwise we had a 3-node, and the two inhabitants, together with k , momentarily form a node with three elements; in sorted order, call them k_1, k_2 and k_3 .
- We now **split** the node, so that k_1 and k_3 form their own individual 2-nodes. The middle key, k_2 is **promoted** to the parent node.
- The promotion may cause the parent node to overflow, in which case **it** gets split the same way. The only time the tree's height changes is when the root overflows.

Example: Build a 2-3 Tree from 9, 5, 8, 3, 2, 4, 7



Example: Build a 2-3 Tree from 9, 5, 8, 3, 2, 4, 7




Exercise: 2-3 Tree Construction: C, O, M, P, U, T, I, N, G


- Build the 2-3 tree that results from inserting these keys, in the given order, into an initially empty tree:

C, O, M, P, U, T, I, N, G

Exercise: 2-3 Tree Construction: C, O, M, P, U, T, I, N, G



Exercise: 2-3 Tree Construction: C, O, M, P, U, T, I, N, G



2-3 Tree Analysis

- Worst case search time results when all nodes are 2-nodes.
- The relation between the number n of nodes and the height h is:

$$n = 1 + 2 + 4 + \dots + 2^h = 2^{h+1} - 1$$

- That is, $\log_2(n + 1) = h + 1$.
- In the best case, all nodes are 3-nodes:

$$n = 2 + 2 \times 3 + 2 \times 3^2 + \dots + 2 \times 3^h = 3^{h+1} - 1$$

- That is, $\log_3(n + 1) = h + 1$.
- Hence we have $\log_3(n + 1) - 1 \leq h \leq \log_2(n + 1) - 1$.
- Useful formula: $\sum_{i=0}^n a^i = \frac{a^{n+1} - 1}{a - 1}$, for $a \neq 1$.

Coming Up Next

- How to buy time, by spending a bit more space (Levitin 7.1 & 7.2).