

David Eccles



INFO90002

Database Systems & Information Modelling

L13 Transactions & Concurrency



- Subject Experience Survey
 - Many changes this summer!
 - Your chance to pay it forward
 - You have benefited from your peer's suggestions
 - Keep it constructive & professional!



<https://ses.unimelb.edu.au/>

We listen.
We act.
YOU benefit.



MELBOURNE

- Hour 1: Transactions & Concurrency
- Hour 2: Storage and Indexes



- Why we need user-defined transactions
- Properties of transactions
- How to use transactions
- Concurrent access to data
- Locking and deadlocking
- Database recovery



- Example “business transactions”:
 - Insert one row in *Order* table, then several in *OrderItem* table
 - Check amount < balance. If so, subtract amount from one row in bank account table, then add amount to another row
 - For all rows in Customer table, send out monthly statements
- Each requires several distinct database operations ...
- e.g. move money from savings account to credit card account
 - Accept inputs from user (via ATM, internet banking or mobile app)
 - Select balance from savings account
 - Is there enough money to withdraw? If so:
 - Update savings account balance = balance – withdrawal
 - Update credit card balance = balance + withdrawal
 - If no errors encountered, end successfully



- A logical unit of work that must either be entirely completed or aborted (indivisible, atomic)
- DML statements are already atomic in MySQL*, SQL Server*
- RDBMS also allows for *user-defined* transactions
- These are a sequence of DML statements, such as
 - a series of UPDATE statements to change values
 - a series of INSERT statements to add rows to tables
 - DELETE statements to remove rows
- Transactions will be treated as atomic
- A successful transaction changes the database from one consistent state to another
 - All data integrity constraints are satisfied

* Database specific, e.g Oracle DML are not atomic transactions by default



- Transactions solve TWO problems:
 1. users need the ability to define a unit of work
 2. concurrent access to data by >1 user or program
- (also acts as an “undo” for manual database manipulation)



MELBOURNE

- Single DML or DDL command (implicit transaction)
 - e.g. Update 700 records,
but database crashes after 200 records processed
 - Restart server: you will find no changes to any records
 - Changes are “all or none”
- Multiple statements (user-defined transaction)
 - **START TRANSACTION;** (or, ‘**BEGIN**’)
 - **SQL statement;**
 - **SQL statement;**
 - **SQL statement;**
 - ...
 - **COMMIT;** (commits the whole transaction)
 - Or **ROLLBACK** (to undo everything)
 - SQL keywords: **BEGIN; COMMIT, ROLLBACK**



- Each transaction consists of several SQL statements, embedded within a larger application program
- Transaction needs to be an indivisible unit of work
 - “Indivisible” means that either the whole job gets done, or none gets done:
 - if an error occurs, we don’t leave the database with the job half done, in an inconsistent state

In the case of an error:

- Any SQL statements already completed must be reversed
- Show an error message to the user
- When ready, the user can try the transaction again
- This is briefly annoying – but inconsistent data is disastrous



- Two scripts `cre_account.sql` and `account.sql` available from the resources page.

```
9   -- Transaction;
10 •  START TRANSACTION; -- An explicit start - but after any commit a NEW transaction begins
11
12   -- Statement 2
13 •  SELECT * FROM ACCOUNT;
14
15   -- (declare a temporary variable amount persistent for this session)
16 •  set @amount = 100;
17
18   -- Statement 3
19
20 •  UPDATE ACCOUNT set balance = balance - @amount where id =1;
21
22   -- Statement 4 confirm deduction from savings but not yet deposited to credit
23 •  SELECT * FROM ACCOUNT;
24
25   -- Statement 5 deposit the amount into the credit account
26 •  UPDATE ACCOUNT set balance = balance + @amount where id = 2;
27
28   -- Statement 6 confirm all changes
29 •  SELECT * FROM ACCOUNT;
30
31   -- Statement 7 EXPLICIT COMMIT;
32 •  COMMIT;
33
34   -- ALL CHANGES PERMANENT CAN NOT BE UNDONE WITH ROLLBACK
35   ...
```



- **Atomicity**
 - A transaction is treated as a single, indivisible, logical unit of work. All operations in a transaction must be completed; if not, then the transaction is aborted
- **Consistency**
 - Constraints that hold before a transaction must also hold after it
 - (multiple users accessing the same data see the same value)
- **Isolation**
 - Changes made during execution of a transaction cannot be seen by other transactions until this one is completed
- **Durability**
 - When a transaction is complete, the changes made to the database are permanent, even if the system fails



- What happens if we have multiple users accessing the database at the same time...
- Concurrent execution of DML against a shared database
- Note that the sharing of data among multiple users is where much of the benefit of databases derives – users communicate and collaborate via shared data
- But what could go wrong?
 - lost updates
 - uncommitted data
 - inconsistent retrievals



MELBOURNE

Alice



Read account
balance
(balance = \$1000)

Withdraw \$100
(balance = \$900)

Write balance
balance = \$900



Bob



Read account
balance
(balance = \$1000)

Withdraw \$800
(balance = \$200)

Write balance
balance = \$200

Balance should be \$100



- Uncommitted data occurs when two transactions execute concurrently and the first is rolled back after the second has already accessed the uncommitted data

Alice



Time

t_1

t_3

t_2

t_5

t_6

Read balance Withdraw \$100 Write balance
(balance = \$200) (balance = \$100) **balance = \$100**

Bob

Read balance Withdraw \$800
(balance = \$1000) (balance = \$200)

Rollback
balance = \$1000

Balance should be \$900



- Occurs when one transaction calculates some aggregate functions over a set of data, while other transactions are updating the data
 - Some data may be read after they are changed and some before they are changed, yielding inconsistent results

| Alice | Bob |
|--|---|
| <pre>SELECT SUM(Salary) FROM Employee;</pre> | <pre>UPDATE Employee SET Salary = Salary * 1.01 WHERE EmpID = 33;</pre> |
| | <pre>UPDATE Employee SET Salary = Salary * 1.01 WHERE EmpID = 44;</pre> |
| (finishes calculating sum) | COMMIT; |



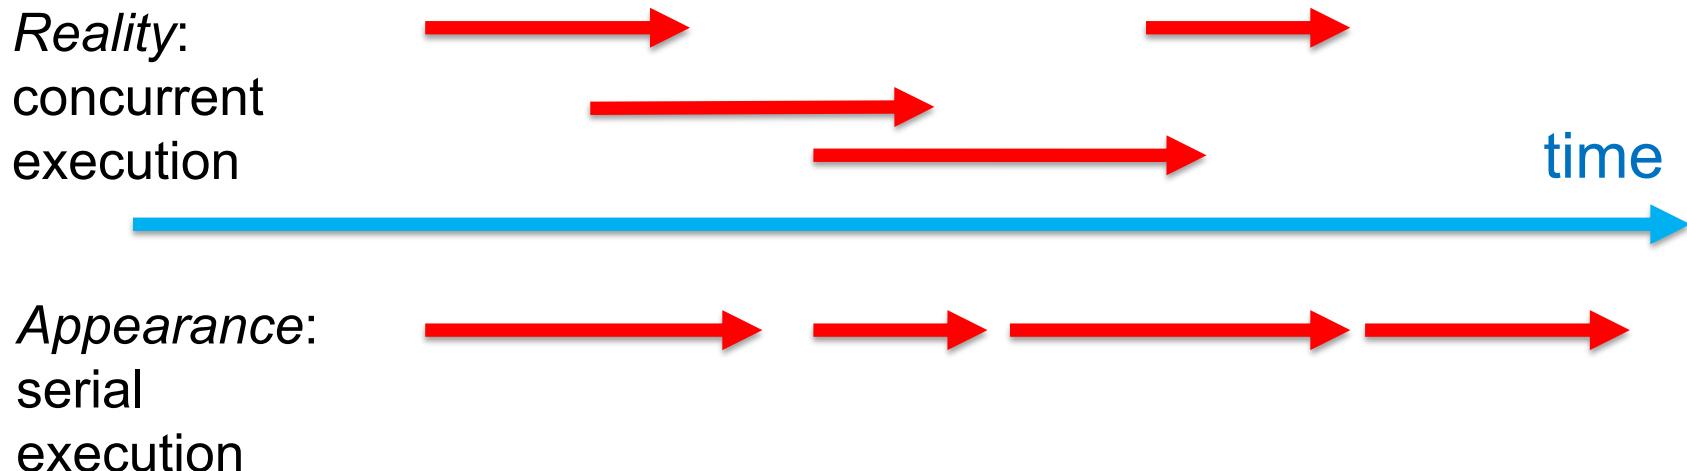
Example: Inconsistent Retrieval

© University of Melbourne

| Time | Trans- action | Action | Value | T1 SUM | Comment |
|------|------------------|---------------------------|--------|---------|--|
| 1 | T1 | Read Salary for EmpID 11 | 10,000 | 10,000 | |
| 2 | T1 | Read Salary for EmpID 22 | 20,000 | 30,000 | |
| 3 | T2 | Read Salary for EmpID 33 | 30,000 | | |
| 4 | T2 | Salary = Salary * 1.01 | | | |
| 5 | T2 | Write Salary for EmpID 33 | 30,300 | | |
| 6 | T1 | Read Salary for EmpID 33 | 30,300 | 60,300 | after update |
| 7 | T1 | Read Salary for EmpID 44 | 40,000 | 100,300 | before update |
| 8 | T2 | Read Salary for EmpID 44 | 40,000 | | |
| 9 | T2 | Salary = Salary * 1.01 | | | we want either before \$210,000 or after \$210,700 |
| 10 | T2 | Write Salary for EmpID 44 | 40,400 | | |
| 11 | T2 | COMMIT | | | |
| 12 | T1 | Read Salary for EmpID 55 | 50,000 | 150,300 | |
| 13 | T1 | Read Salary for EmpID 66 | 60,000 | 210,300 | |



- Transactions ideally are “serializable”
 - Multiple, concurrent transactions appear as if they were executed one after another
 - Ensures that the concurrent execution of several transactions yields consistent results



but true serial execution (i.e. no concurrency) is very expensive!



- To achieve efficient execution of transactions, the DBMS creates a schedule of read and write operations for concurrent transactions
- Interleaves the execution of operations, based on concurrency control algorithms such as locking and time stamping
- Several methods of concurrency control
 - *Locking* is the main method used
 - Alternate methods
 - Time Stamping
 - Optimistic Methods



- Lock
 - Guarantees exclusive use of a data item to a current transaction
 - T1 acquires a lock prior to data access; the lock is released when the transaction is complete
 - T2 does not have access to data item currently being used by T1
 - T2 has to wait until T1 releases the lock
 - Required to prevent another transaction from reading inconsistent data
- Lock manager
 - Responsible for assigning and policing the locks used by the transactions
- Question: at what granularity should we apply locks?



- Database-level lock
 - Entire database is locked
 - Good for batch processing but unsuitable for multi-user DBMSs
 - T1 and T2 can not access the same database concurrently even if they use different tables
 - (SQLite, Access)
- Table-level lock
 - Entire table is locked - as above but not quite as bad
 - T1 and T2 can access the same database concurrently as long as they use different tables
 - Can cause bottlenecks, even if transactions want to access different parts of the table and would not interfere with each other
 - Not suitable for highly multi-user DBMSs



- **Page-level lock**
 - An entire disk page is locked (a table can span several pages and each page can contain several rows of one or more tables)
 - Not commonly used now
- **Row-level lock**
 - Allows concurrent transactions to access different rows of the same table, even if the rows are located on the same page
 - Improves data availability but with high overhead (each row has a lock that must be read and written to)
 - Currently the most popular approach (MySQL, Oracle)
- **Field-level lock**
 - Allows concurrent transactions to access the same row, as long as they access different attributes within that row
 - Most flexible lock but requires an extremely high level of overhead
 - Not commonly used



- Binary Locks
 - has only two states: locked (1) or unlocked (0)
 - eliminates “Lost Update” problem
 - the lock is not released until the statement is completed
 - considered too restrictive to yield optimal concurrency, as it locks even for two READs (when no update is being done)
- The alternative is to allow both *Exclusive* and *Shared* locks
 - often called Read and Write locks
- Writers never block Readers
 - INSERT UPDATE DELETE should never block a SELECT
- Readers never block Writers
 - SELECT should never block INSERT UPDATE DELETE
- Read Consistent; Read Committed are 2 options



- **Exclusive lock**
 - access is reserved for the transaction that locked the object
 - must be used when transaction intends to WRITE
 - granted if and only if no other locks are held on the data item
 - in MySQL: “select ... for update”
- **Shared lock**
 - other transactions are also granted Read access
 - issued when a transaction wants to READ data, and no Exclusive lock is held on that data item
 - multiple transactions can each have a shared lock on the same data item if they are all just reading it
 - in MySQL: “select ... lock in share mode”



- Two Rules
 - 1. If a transaction T wants to read an object O , it first requests a *shared* lock on the object and then obtains an *exclusive* lock
 - 2. All locks held by transaction T are released when the transaction is completed.
- In effect the locking protocol only allows 'safe' interleaving of transactions T_i . If two transactions are accessing completely different parts of the dbms they concurrently obtain the locks and proceed.
- If two transactions T_1, T_2 are accessing the same object and one wants to modify it, their actions are ordered serially – all actions of one transaction T_1 must complete before the other transaction T_2 can proceed.



Demo: Concurrent execution

- I will demonstrate locking and concurrency with transactions.

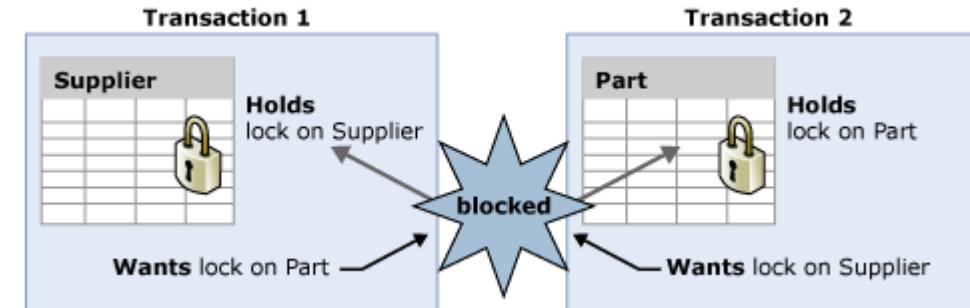
```
H:>>
H:>>mysql -u root test
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 21
Server version: 5.6.21 MySQL Community Server (GPL)

Copyright (c) 2000, 2014, Oracle and/or its affiliates. All
Oracle is a registered trademark of Oracle Corporation and/
affiliates. Other names may be trademarks of their respecti
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the curre
mysql> select * from Account;
```



- Condition that occurs when two transactions wait for each other to unlock data
 - T1 locks data item X, then wants Y
 - T2 locks data item Y, then wants X
 - each waits to get a data item which the other transaction is already holding
 - could wait forever if not dealt with
- Only happens with **exclusive** locks
- Deadlocks are dealt with by:
 - prevention
 - detection
 - (we won't go into the details of how in this course)





```
deccles2 — mysql -u root -p — 80x24
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
Database changed
mysql> SELECT * FROM DEPT;
+-----+-----+-----+
| deptno | name      | location   |
+-----+-----+-----+
|    10  | ACCOUNTING | NEW YORK  |
|    20  | RESEARCH   | DALLAS    |
|    30  | SALES      | CHICAGO   |
|    40  | OPERATIONS | BOSTON    |
+-----+-----+-----+
4 rows in set (0.00 sec)

1 mysql> UPDATE DEPT set location = 'SYDNEY' where deptno = 30;
Query OK, 1 row affected (0.08 sec)
Rows matched: 1  Changed: 1  Warnings: 0

3 mysql> UPDATE DEPT set location = 'MELBOURNE' where deptno = 20;
Query OK, 1 row affected (13.07 sec)
Rows matched: 1  Changed: 1  Warnings: 0

mysql> [REDACTED]

deccles2 — mysql -u root -p — 71x24
mysql> USE SCOTT;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
Database changed
mysql> SELECT * FROM DEPT;
+-----+-----+-----+
| deptno | name      | location   |
+-----+-----+-----+
|    10  | ACCOUNTING | NEW YORK  |
|    20  | RESEARCH   | DALLAS    |
|    30  | SALES      | CHICAGO   |
|    40  | OPERATIONS | BOSTON    |
+-----+-----+-----+
4 rows in set (0.00 sec)

2 mysql> UPDATE DEPT set location = 'AUCKLAND' WHERE deptno = 20;
Query OK, 1 row affected (0.00 sec)
Rows matched: 1  Changed: 1  Warnings: 0

4 mysql> Update dept set location = 'WELLINGTON' where deptno = 30;
ERROR 1213 (40001): Deadlock found when trying to get lock; try restarting transaction
mysql>
```

- Two separate sessions
- In order
- Tx1 Update row 3 (Green)
- Tx2 Update row 2 (White)
- Tx3 Update row 2 (Green)
- Tx4 Update row 3 (White)
- Note: Only the session which detects the deadlock rolls back the transaction.
The Green session still holds locks on row 2 and 3



- Isolation levels provide the level any given transaction is exposed to the actions of other executing transactions
- Isolation levels are *READ UNCOMMITTED*, *READ COMMITTED*, *REPEATABLE READ* and *SERIALIZABLE*
- *Serializable* is the highest degree of isolation
 - Locks are obtained for both reads (R) and writes (W)
 - Transaction T only reads committed transactions
- *Repeatable Read*
 - T only reads changes made by committed transactions and no value read or written by T can be changed by any other transaction until T completes.



- ***READ COMMITTED***
 - Only reads committed changes by other transactions
 - No value written by T is changed by another transaction until T completes (commits)
 - However a value read by T can be modified by another transactions while T is in progress.
 - Read committed holds exclusive locks (x) before writing objects and shared locks (s) before reading objects.
- ***READ UNCOMMITTED***
 - Does not obtain shared locks before reading
 - Prohibits any transaction from making changes (i.e. write)
 - Access mode is *read only*
- **SET TRANSACTION ISOLATION LEVEL <level>;**



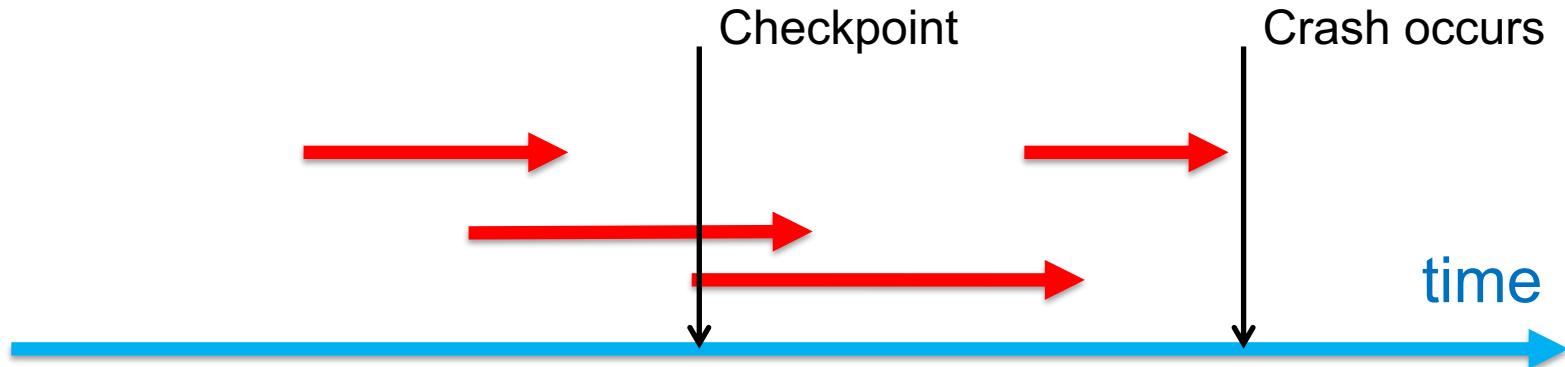
- **Timestamp**
 - Assigns a global unique timestamp to each transaction
 - Each data item accessed by the transaction gets the timestamp
 - Thus for every data item, the DBMS knows which transaction performed the last read or write on it
 - When a transaction wants to read or write, the DBMS compares its timestamp with the timestamps already attached to the item and decides whether to allow access
- **Optimistic**
 - Based on the assumption that the majority of database operations do not conflict
 - Transaction is executed without restrictions or checking
 - Then when it is ready to commit, the DBMS checks whether it any of the data it read has been altered – if so, rollback



- Want to restore database to a previous consistent state
- If transaction cannot be completed, it must be aborted and any changes rolled back
- To enable this, DBMS tracks all updates to data
- This *transaction log* contains:
 - a record for the beginning of the transaction
 - for each SQL statement
 - operation being performed (update, delete, insert)
 - objects affected by the transaction
 - “before” and “after” values for updated fields
 - pointers to previous and next transaction log entries
 - the ending (COMMIT) of the transaction



- Also provides the ability to restore a corrupted database
- If a system failure occurs, the DBMS will examine the log for all uncommitted or incomplete transactions and it will restore the database to a previous state





Example transaction log

| TRL ID | TRX NUM | PREV PTR | NEXT PTR | OPERATION | TABLE | ROW ID | ATTRIBUTE | BEFORE VALUE | AFTER VALUE |
|--------|---------|----------|----------|------------|-------------------------|----------|--------------|--------------|-------------------------|
| 341 | 101 | Null | 352 | START | ****Start Transaction | | | | |
| 352 | 101 | 341 | 363 | UPDATE | PRODUCT | 54778-2T | PROD_QOH | 45 | 43 |
| 363 | 101 | 352 | 365 | UPDATE | CUSTOMER | 10011 | CUST_BALANCE | 615.73 | 675.62 |
| 365 | 101 | 363 | Null | COMMIT | **** End of Transaction | | | | |
| 397 | 106 | Null | 405 | START | ****Start Transaction | | | | |
| 405 | 106 | 397 | 415 | INSERT | INVOICE | 1009 | | | 1009,10016, ... |
| 415 | 106 | 405 | 419 | INSERT | LINE | 1009,1 | | | 1009,1, 89-WRE-Q,1, ... |
| 419 | 106 | 415 | 427 | UPDATE | PRODUCT | 89-WRE-Q | PROD_QOH | 12 | 11 |
| 423 | | | | CHECKPOINT | | | | | |
| 427 | 106 | 419 | 431 | UPDATE | CUSTOMER | 10016 | CUST_BALANCE | 0.00 | 277.55 |
| 431 | 106 | 427 | 457 | INSERT | ACCT_TRANSACTION | 10007 | | | 1007,18-JAN-2004, ... |
| 457 | 106 | 431 | Null | COMMIT | **** End of Transaction | | | | |
| 521 | 155 | Null | 525 | START | ****Start Transaction | | | | |
| 525 | 155 | 521 | 528 | UPDATE | PRODUCT | 2232/QWE | PROD_QOH | 6 | 26 |
| 528 | 155 | 525 | Null | COMMIT | **** End of Transaction | | | | |

***** C *R*A* S* H *****



- Why we need user-defined transactions
- Properties of transactions
 - ACID
- How to use transactions
 - BEGIN TRANSACTION; START; COMMIT; ROLLBACK;
- Concurrent access to data
 - Concurrent access strategies
- Locking and deadlocking
 - Types of Locks (Binary; Shared)
- Database recovery
 - Fundamentals of transaction recovery

* All material is examinable – these are the suggested key skills you would need to demonstrate in an exam scenario



- Storage and Indexing



David Eccles

THE UNIVERSITY OF
MELBOURNE

L13. Transactions & Concurrency

INFO90002 Database Systems & Information Modelling